

# Mining Data Streams

Liyao Xiang  
[xiangliyao.cn](http://xiangliyao.cn)

Shanghai Jiao Tong University

- What is streaming data and its application?

# Data Streams

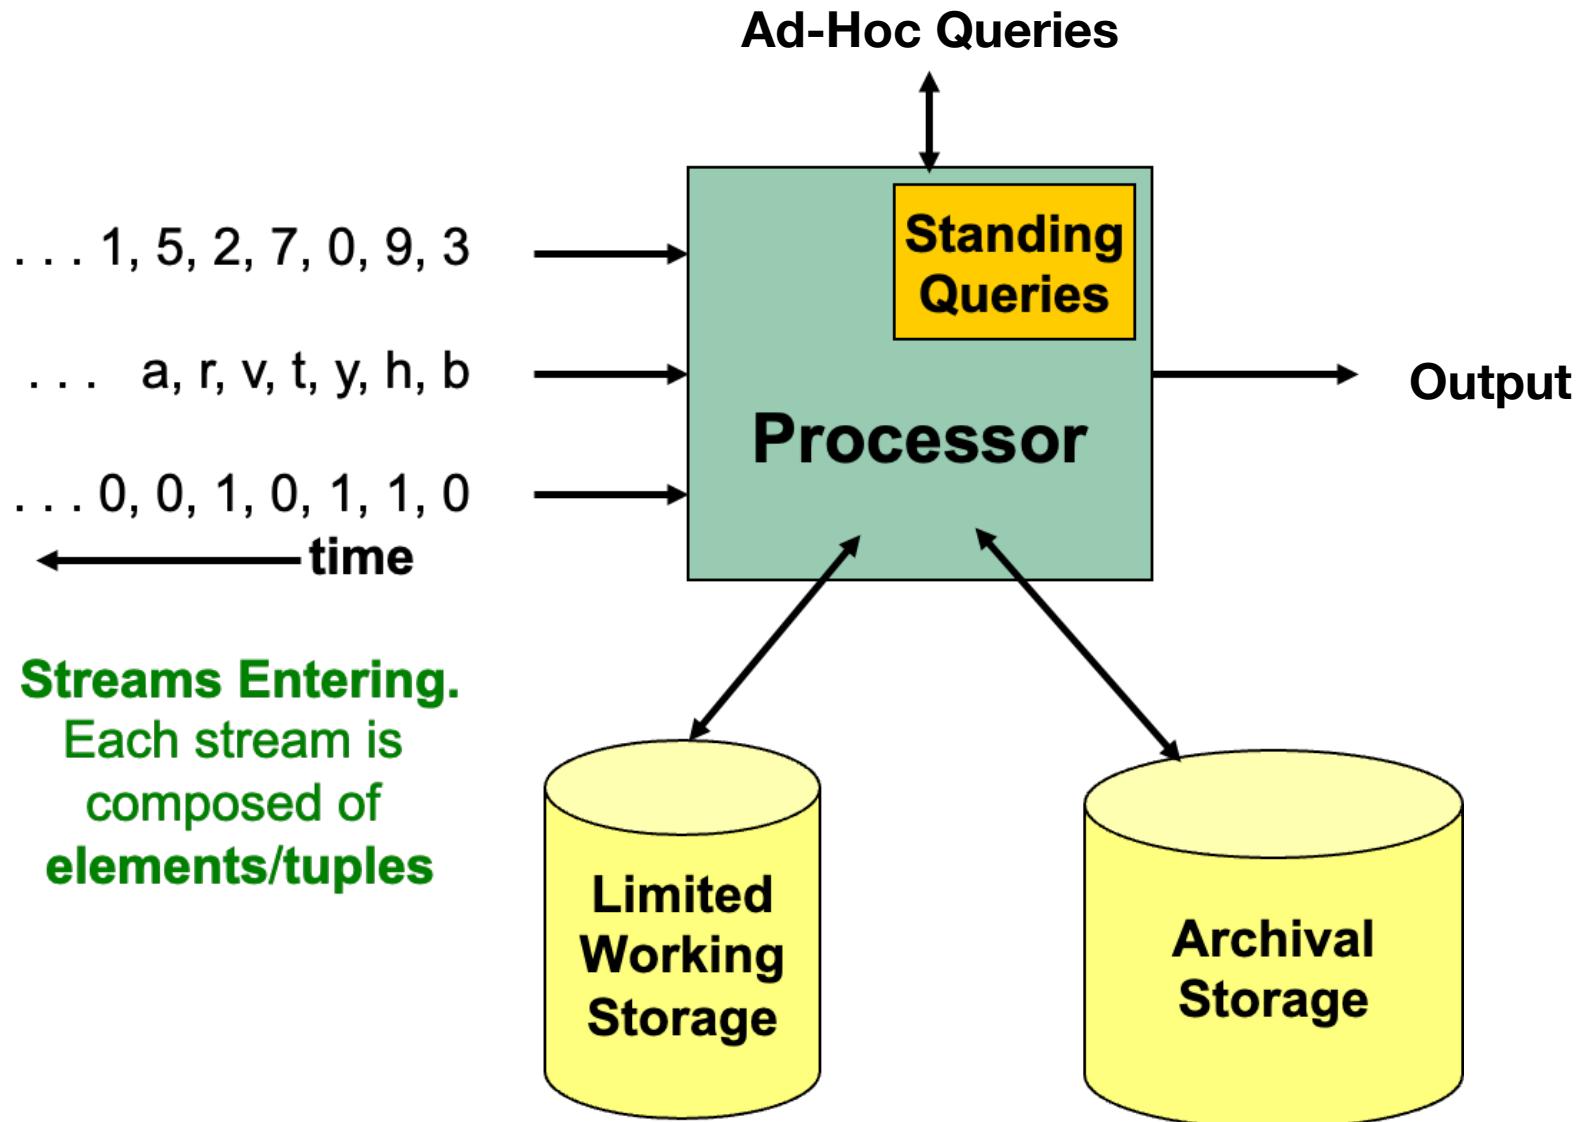
- In many data mining situations, we do not know the entire data set in advance
- **Stream Management** is important when the input rate is controlled **externally**:
  - Baidu queries
  - Click-through-rate on bilibili
  - Weibo status/blog updates
- We can think of the data as **infinite** and **non-stationary** (the distribution changes over time)

# The Stream Model

- Input elements enter at a rapid rate, at one or more input ports (i.e., streams)
  - We call elements of the stream **tuples**
- The system **cannot** store the entire stream accessibly
- Problem: How do we make critical calculations about the stream using **a limited amount of (secondary) memory?**

**Creating summary!**

# General Stream Processing Model



# Example: SGD is a Streaming Alg.

- Stochastic Gradient Descent (SGD) is an example of a stream alg.
- In Machine Learning we call this: **Online Learning**
  - Allows for modeling problems where we have a **continuous stream of data**
  - We want an algorithm to learn from it and **slowly adapt** to the changes in data
- Idea: Do **slow updates** to the model
  - SGD makes small updates
  - So: First train the classifier on training data
  - Then: For every example from the stream, we slightly update the model (using small learning rate)

# What are the Queries?

- What is the data distribution of the stream?
- How many items featured by XXX in the last K elements of the stream?
- What are the items with property YY from the stream?
- How many distinctive items in the last K elements of the stream?
- What are the statistics of the last K elements?
- What is the most frequently appearing items in the stream?
- ...

# Applications

- Mining query streams
  - Baidu wants to know what queries are **more frequent** today than yesterday
- Mining click streams
  - Bilibili wants to know which of the video clips are **getting an unusual high number** of hits in the past hour
- Mining social network news feeds
  - Weibo looks for **trending topics**

# Applications

- Sensor networks where many sensors' feeding data into a central controller
  - Want to detect unusual events
- IP packets monitored at a switch
  - Gather information for optimal routing
  - Detect denial-of-service attacks

- How do we sample data from a stream?

# Sampling from a Data Stream

- Since we **can not store the entire stream**, one obvious approach is to store **a sample**
- How do we sample from the stream just **like we have seen the entire stream?**
  - At any “time”  $k$  we would like a random sample of  $s$  elements
  - **What is the property of the sample we want to maintain?**
    - For all time steps  $k$ , each of  $k$  elements seen so far has equal prob. of being sampled
- Two different problems:
  1. Sample a fixed proportion
  2. Sample a fixed-sized set

# Sampling a Fixed Proportion

- Problem 1: Sample a **fixed proportion** of elements in the stream (say 1 in 10)
- Scenario: Search engine query stream
  - Stream of tuples: (user, query, time)
  - Answer questions such as: How often did a user run the same query through the search engine in a single day?
  - Have space to store 1/10th of query stream
- Naïve solution:
  - Generate a random integer in [0..9] for each query
  - Store the query if the integer is 0, otherwise discard

# Problem with Naive Approach

- Simple question: What fraction of queries by an average search engine user are duplicates?
  - Suppose each user issues  $x$  queries once and  $d$  queries twice (total of  $x+2d$  queries)
- Proposed solution: We keep 10% of the queries
  - Sample will contain  $x/10$  of the singleton queries and  $2d/10$  of the duplicate queries at least once
  - But only  $d/100$  pairs of duplicates
    - $d/100 = 1/10 \cdot 1/10 \cdot d$
    - Of  $d$  “duplicates”  $18d/100$  appear exactly once
      - $18d/100 = ((1/10 \cdot 9/10) + (9/10 \cdot 1/10)) \cdot d$
  - What fraction of queries by an average search engine user are duplicates by the **sample-based** approach?

# Sampling Users

- How about sampling users?
  - Pick 1/10th of **users** and take all their searches in the sample
  - Use a hash function that hashes the user name or user id uniformly into 10 buckets
- What fraction of queries by an average search engine user are duplicates by the **user-based** approach?

# Generalized Solution

- Stream of tuples with keys:
  - Key is some subset of each tuple's components
    - e.g., tuple is (user, search query, time); key is **user**
  - Choice of key depends on application
- To get a sample of  $a/b$  fraction of the stream:
  - Hash each tuple's key uniformly into **b** buckets
  - Pick the tuple if its hash value is at most **a**



Hash table with **b** buckets, pick the tuple if its hash value is at most **a**.  
How to generate a 30% sample? Hash into  $b=10$  buckets, take the tuple if it  
hashes to one of the first 3 buckets

# Sampling a fixed-size sample

- Problem 2: Fixed-size sample
  - As the stream grows, the sample is of fixed size
- Suppose we need to maintain a random sample  $S$  of size exactly  $s$  tuples
  - E.g., main memory size constraint is  $s$
- Why? Don't know length of stream in advance
- Suppose at time  $n$  we have seen  $n$  items
  - Each item is in the sample  $S$  with equal probability  $s/n$

**How to think about the problem: say  $s = 2$**

**Stream:** a x c y z k c d e g...

At  $n=5$ , each of the first 5 tuples is included in the sample  $S$  with equal prob.

At  $n=7$ , each of the first 7 tuples is included in the sample  $S$  with equal prob.

An impractical solution would be to store all the  $n$  tuples seen so far and out of them pick  $s$  at random

# Solution: Fixed Size Sample

- **Algorithm (a.k.a. Reservoir Sampling)**
  - Store all the first  $s$  elements of the stream to  $\mathbf{S}$
  - Suppose we have seen  $n-1$  elements, and now the  $n^{\text{th}}$  element arrives ( $n > s$ )
    - With probability  $s/n$ , keep the  $n^{\text{th}}$  element, else discard it
    - If we picked the  $n^{\text{th}}$  element, then it replaces one of the  $s$  elements in the sample  $\mathbf{S}$ , picked uniformly at random
- Claim: This algorithm maintains a sample  $\mathbf{S}$  with the desired property:
  - After  $n$  elements, the sample contains each element seen so far with probability  $s/n$

# Proof

- We prove this by induction:
  - Assume that after  $n$  elements, the sample contains each element seen so far with probability  $s/n$
  - We need to show that after seeing element  $n+1$  the sample maintains the property
    - Sample contains each element seen so far with probability  $s/(n+1)$
- Base case:
  - After we see  $n=s$  elements the sample  $\mathbf{S}$  has the desired property
    - Each out of  $n=s$  elements is in the sample with probability  $s/s = 1$

# Proof

- **Inductive hypothesis:** After  $n$  elements, the sample  $\mathbf{S}$  contains each element seen so far with prob.  $s/n$
- Now element  $n+1$  arrives
- **Inductive step:** For elements already in  $\mathbf{S}$ , probability that the algorithm keeps it in  $\mathbf{S}$  is:

$$\left(1 - \frac{s}{n+1}\right) + \left(\frac{s}{n+1}\right)\left(\frac{s-1}{s}\right) = \frac{n}{n+1}$$

Element  $n+1$  discarded      Element  $n+1$  not discarded      Element in the sample not picked

- So, at time  $n$ , tuples in  $\mathbf{S}$  were there with prob.  $s/n$
- Time  $n \rightarrow n+1$ , tuple stayed in  $\mathbf{S}$  with prob.  $n/(n+1)$
- So prob. tuple is in  $\mathbf{S}$  at time  $n+1$  =  $\frac{s}{n} \cdot \frac{n}{n+1} = \frac{s}{n+1}$

- How do we query over a sliding window?

# Sliding Windows

- A useful model of stream processing is that queries are about a **window** of length **N** – the **N** most recent elements received
- **Interesting case:** **N** is so large that the data cannot be stored in memory, or even on disk
  - Or, there are so many streams that windows for all cannot be stored
- Example:
  - For every product **X** we keep 0/1 stream of whether that product was sold in **each** transaction
  - We want to answer queries: how many times have we sold **X** in the last **k** sales (cannot know **k** in advance)

# Sliding Window: 1 Stream

- Sliding window on a single stream: **N = 6**

q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

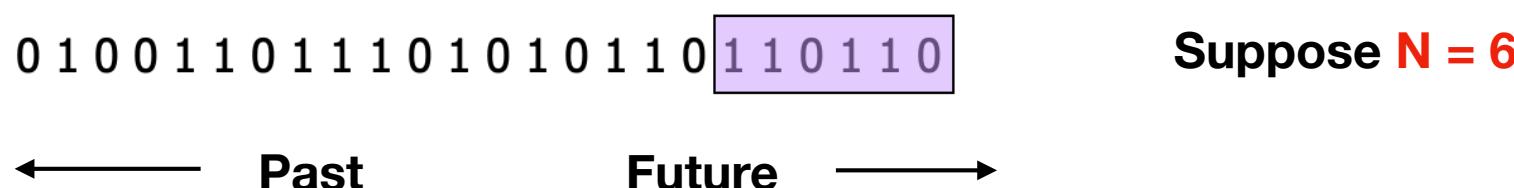
q w e r t y u i o p a s d f g h j k l z x c v b n m

q w e r t y u i o p a s d f g h j k l z x c v b n m

← Past Future →

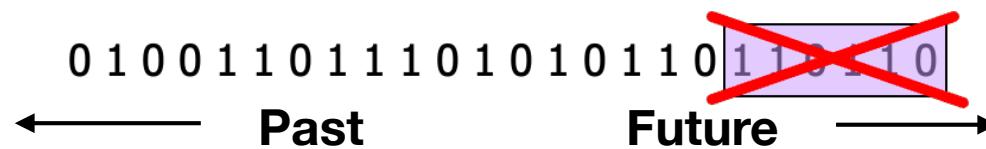
# Counting Bits

- Problem:
  - Given a stream of **0s** and **1s**
  - Be prepared to answer queries of the form:
    - How many 1s are in the last **k** bits? where  $k \leq N$
- Obvious solution:
- Store the most recent **N** bits
  - When new bit comes in, discard the **N+1<sup>th</sup>** bit



# Counting Bits

- You can not get an **exact** answer without storing the **entire** window
- Real Problem: **What if we cannot afford to store N bits?**
  - E.g., we're processing 1 billion streams and  $N = 1$  billion

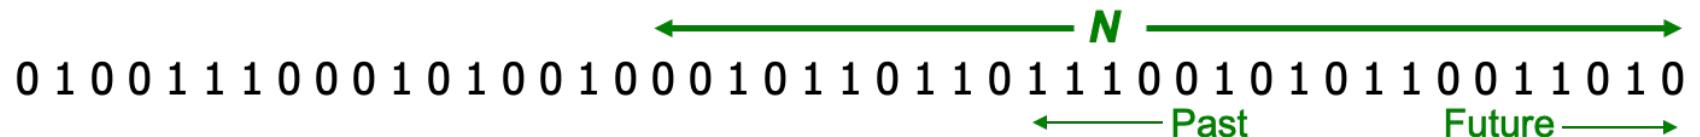


- But we are happy with an **approximate** answer

# An Attempt: Simple Solution

- Q: How many 1s are in the last  $k$  bits?
- A **simple** solution that does not really solve our problem:

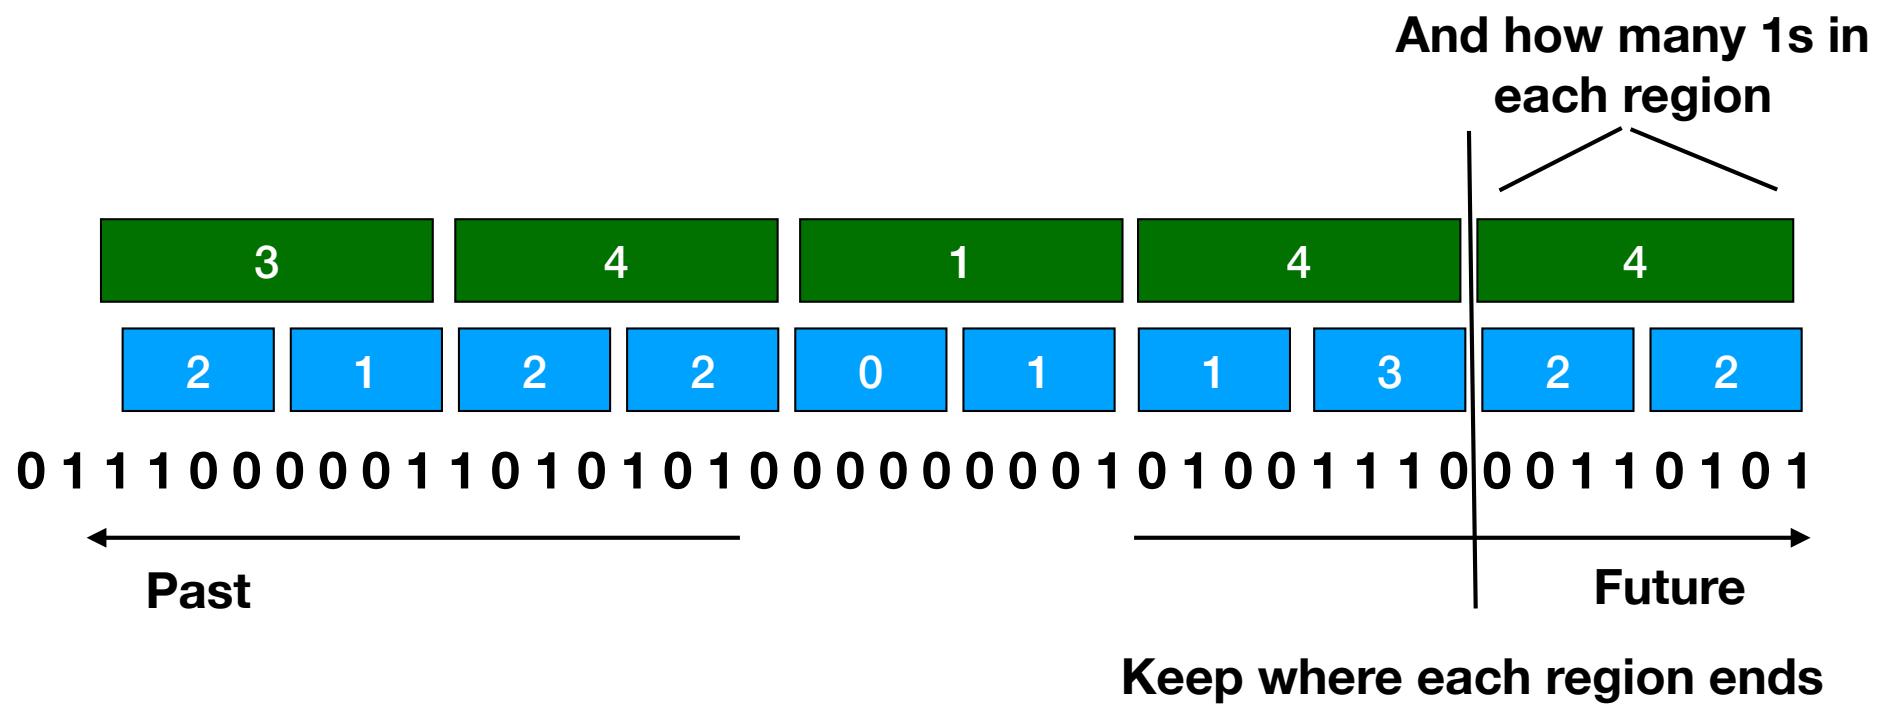
Uniformity assumption



- Maintain 2 counters:
  - $S$ : number of 1s from the beginning of the stream
  - $Z$ : number of 0s from the beginning of the stream
- How many 1s are in the last  $k$  bits?  $k \cdot S / (S+Z)$
- But, what if stream is **non-uniform**?
  - What if distribution changes over time?

# Idea: Uniform Windows

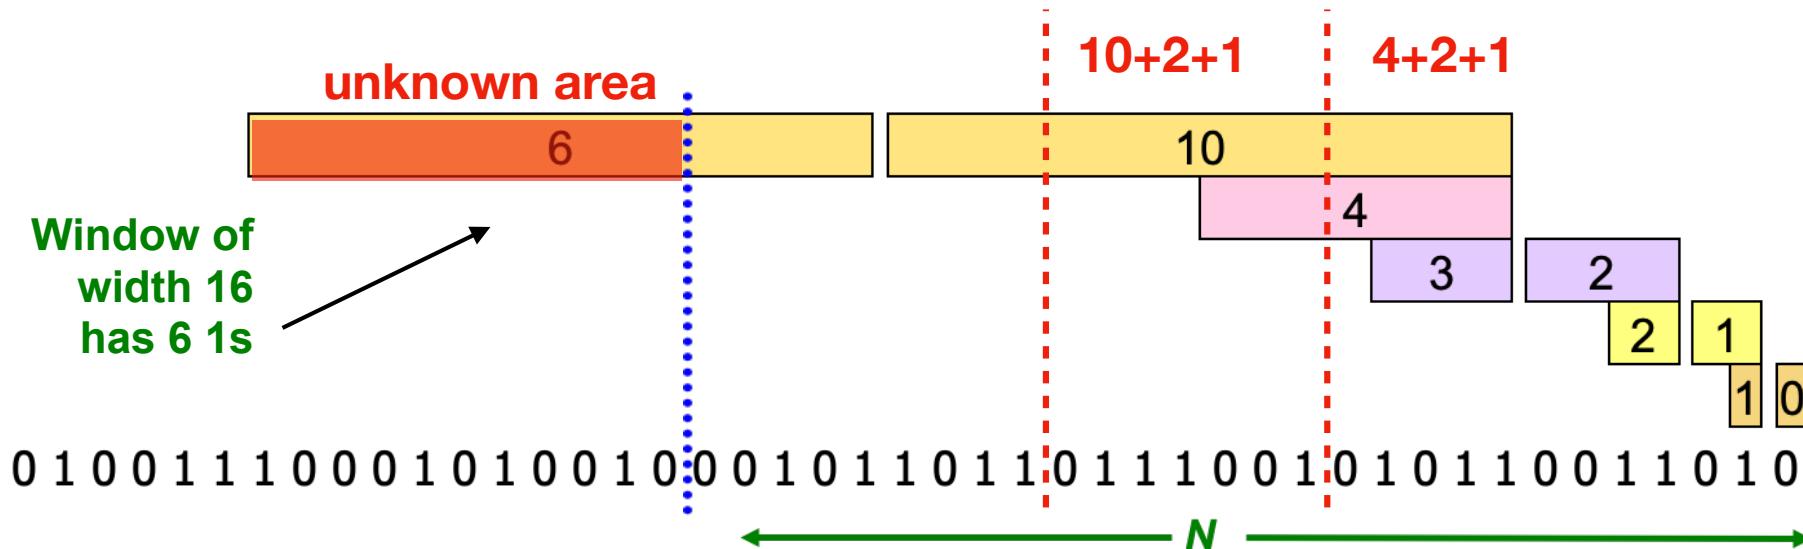
- Summarize **uniformly-sized** regions of the stream, looking backward



- But we usually weigh more on the recent data

# Idea: Exponential Windows

- Summarize **exponentially increasing** regions of the stream, looking backward
  - Drop small regions if they begin at the same point as a larger region



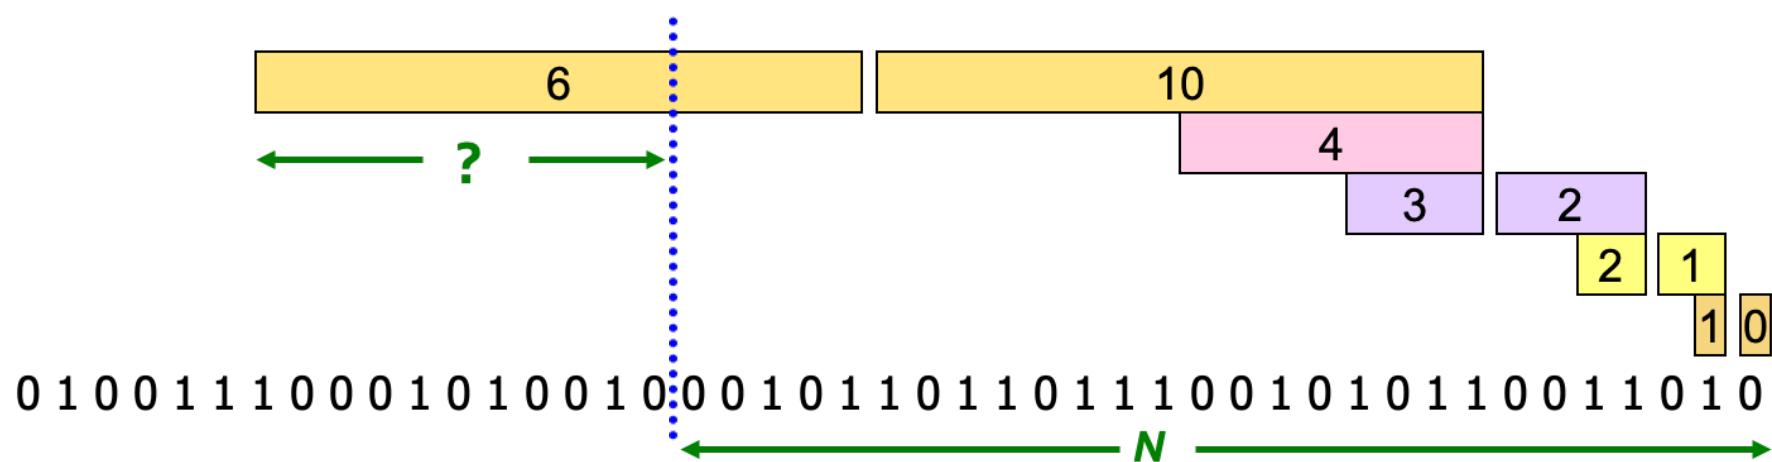
We can reconstruct the count of the last  $N$  bits, except we are not sure how many of the last **6 1s** are included in the  $N$

# Advantages

- Stores only  $O(\log^2 N)$  bits
  - $O(\log N)$  counts of  $\log_2 N$  bits each
- Easy update as more bits enter
- Error in count no greater than the number of 1s in the “**unknown**” area

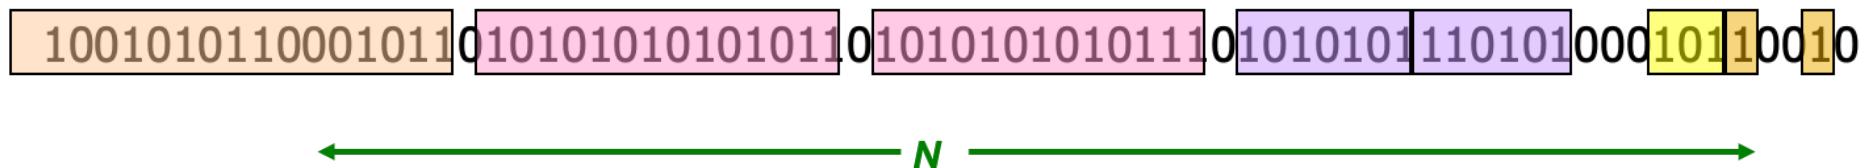
# Drawbacks

- As long as the 1s are fairly evenly distributed, the error due to the unknown region is small – **no more than 50%**
  - But it could be that all the 1s are in the unknown area at the end
  - In that case, the error is unbounded!



# Fixup: DGIM Method

- Idea: Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
  - Let the block **sizes** (number of **1s**) increase exponentially
  - When there are few **1s** in the window, block sizes stay small, so errors are small

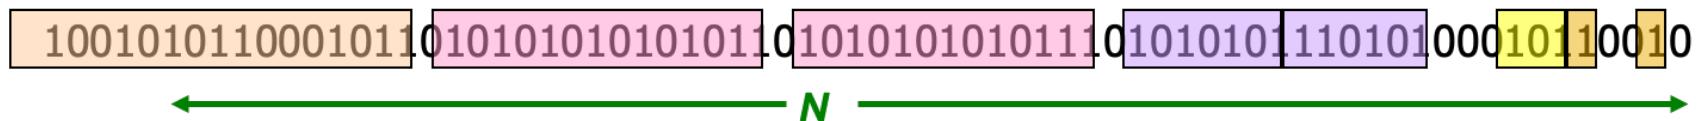


# DGIM: Timestamps

- Each bit in the stream has a timestamp, starting 1, 2, ...
- Record timestamps modulo  $N$  (the window size), so we can represent any relevant timestamp in  $O(\log_2 N)$  bits

# DGIM: Buckets

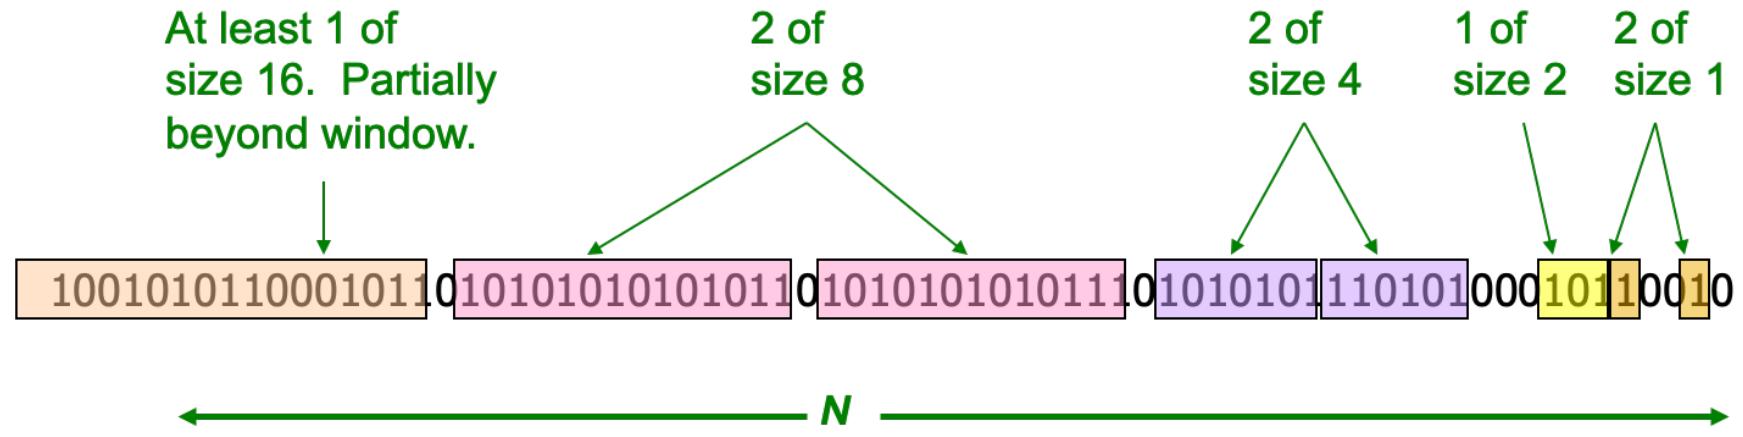
- A bucket in the DGIM method is a record consisting of:
  - A. The timestamp of its end expressed in **[ $O(\log N)$  bits]**
  - B. The number of 1s between its beginning and end expressed in **[ $O(\log N - 1) = O(\log N)$  bits]**
- Constraint on buckets: Number of **1s** must be a power of **2**



# Representing a Stream by Buckets

- Either **one** or **two** buckets with the same **power-of-2 number of 1s**
- Buckets do not overlap in timestamps
- Buckets are sorted by size
  - Earlier buckets are not smaller than later buckets
- Buckets disappear when their end-time is  $> N$  time units in the past

# Example: Bucketized Stream



- Three properties of buckets that are maintained:
  - Either **one** or **two** buckets with the same **power-of-2 number** of 1s
  - Buckets do not overlap in timestamps
  - Buckets are sorted by size

# Updating Buckets

- When a new bit comes in, drop the last (oldest) bucket if its end-time is prior to **N** time units before the current time
- **2 cases:** Current bit is **0** or **1**
- If the current bit is 0: no other changes are needed
- If the current bit is 1:
  1. Create a new bucket of size 1, for just this bit
    - End timestamp = current time
  2. If there are now **three buckets of size 1**, combine the oldest two into **a bucket of size 2**
  3. If there are now **three buckets of size 2**, combine the oldest two into **a bucket of size 4**
  4. And so on ...

# Example: Updating Buckets

**Current state of the stream:**

0010101100010110|101010101010110|10101010101110|1010101110100|10110010

**Bit of value 1 arrives**

0010101100010110|101010101010110|10101010101110|1010101110100|101100101

**Two orange buckets get merged into a yellow bucket**

0010101100010110|101010101010110|10101010101110|1010101110100|101100101

**Next bit 1 arrives, new orange bucket is created, then 0 comes, then 1:**

0101100010110|101010101010110|10101010101110|1010101110100|101100101101

**Buckets get merged...**

0101100010110|101010101010110|10101010101110|1010101110100|101100101101

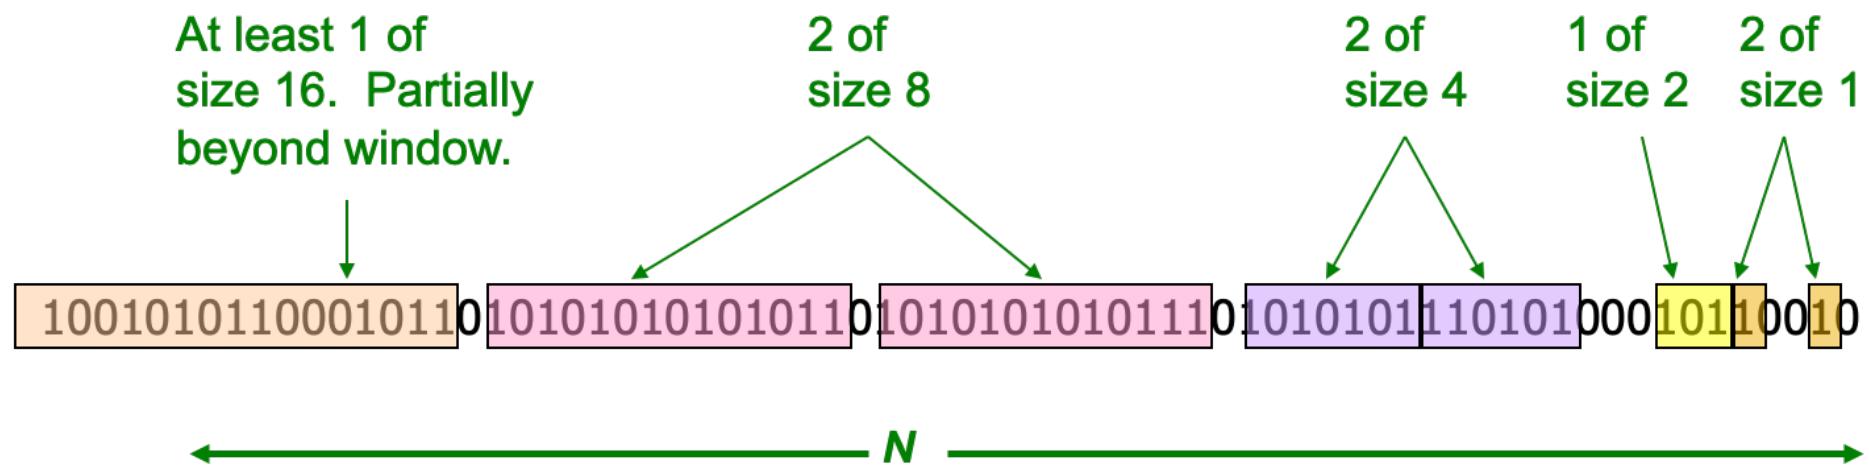
**State of the buckets after merging**

0101100010110|10101010101101010101110|101010111010100|101100101101

# How to Query?

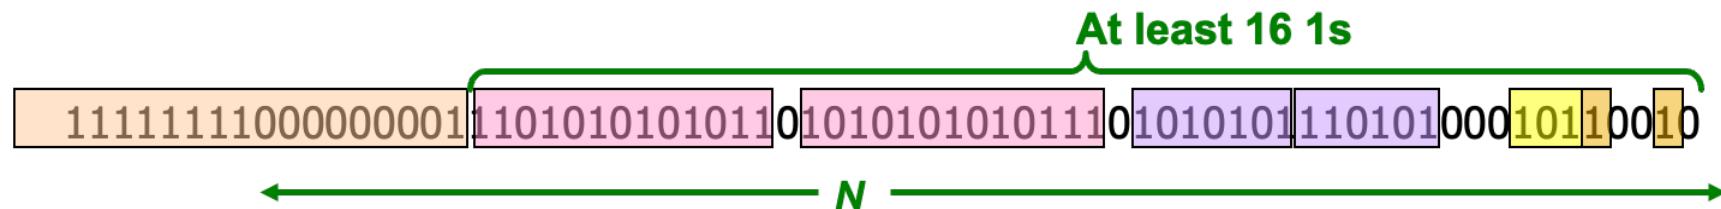
- To estimate the number of **1s** in the most recent **N** bits:
  1. Sum the sizes of all buckets but the last (note “size” means the number of 1s in the bucket)
  2. Add half the size of the last bucket
- Remember: We do not know how many 1s of the last bucket are still within the wanted window

# Example: Bucketized Stream



# Error Bound: Proof

- Why is **error 50%**? Let's prove it!
- Suppose the last bucket has size  $2^r$
- Then by assuming  $2^{r-1}$  (i.e., half) of its 1s are still within the window, we make an error of at most  $2^{r-1}$
- Since there is at least one bucket of each of the sizes less than  $2^r$ , the true sum is at least  $1 + 2 + 4 + \dots + 2^{r-1} = 2^r - 1$
- Thus error is at most 50%



# Further Reducing the Error

- Instead of maintaining **1 or 2** of each size bucket, we allow either **r-1 or r** buckets ( $r > 2$ ) except for the largest-sized buckets
  - We can have any number between **1** and **r** of the largest-sized buckets
- Error is at most **O(1/r)**
- By picking **r** appropriately, we can tradeoff between number of bits we store and the error

# Can We Handle Stream of Integers?

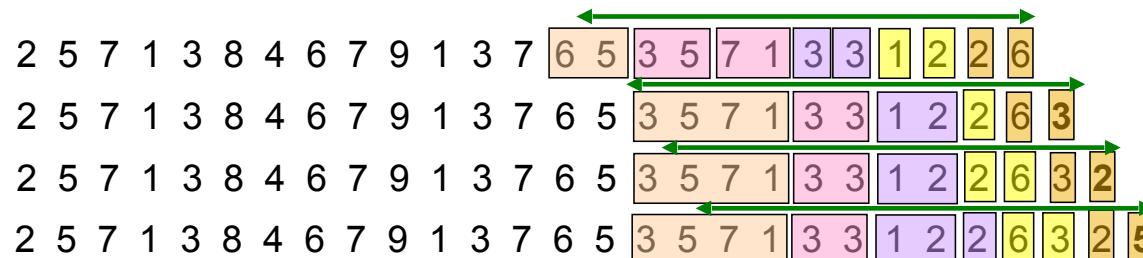
- Stream of positive integers
- We want the sum of the last **k** elements
  - E.g.: Avg. price of last **k** sales
- A naive solution:
  1. If you know all have at most **m** bits
    - Treat **m** bits of each integer as a separate stream
    - Use DGIM to count **1s** in each integer
    - The sum =  $\sum_{i=0}^{m-1} c_i 2^i$       **c<sub>i</sub>** ...estimated count for i-th bit  
**2<sup>i</sup>:** weigh each count differently

# Can We Handle Stream of Integers?

- Stream of positive integers
- We want the sum of the last **k** elements
  - E.g.: Avg. price of last **k** sales
- Another solution:
  2. Use buckets to keep partial sums
    - Sum of elements in size **b** bucket is at most  $2^b$

Idea: Sum in each bucket is at most  $2^b$   
(unless bucket has only 1 integer)  
Bucket sizes:

16 8 4 2 1

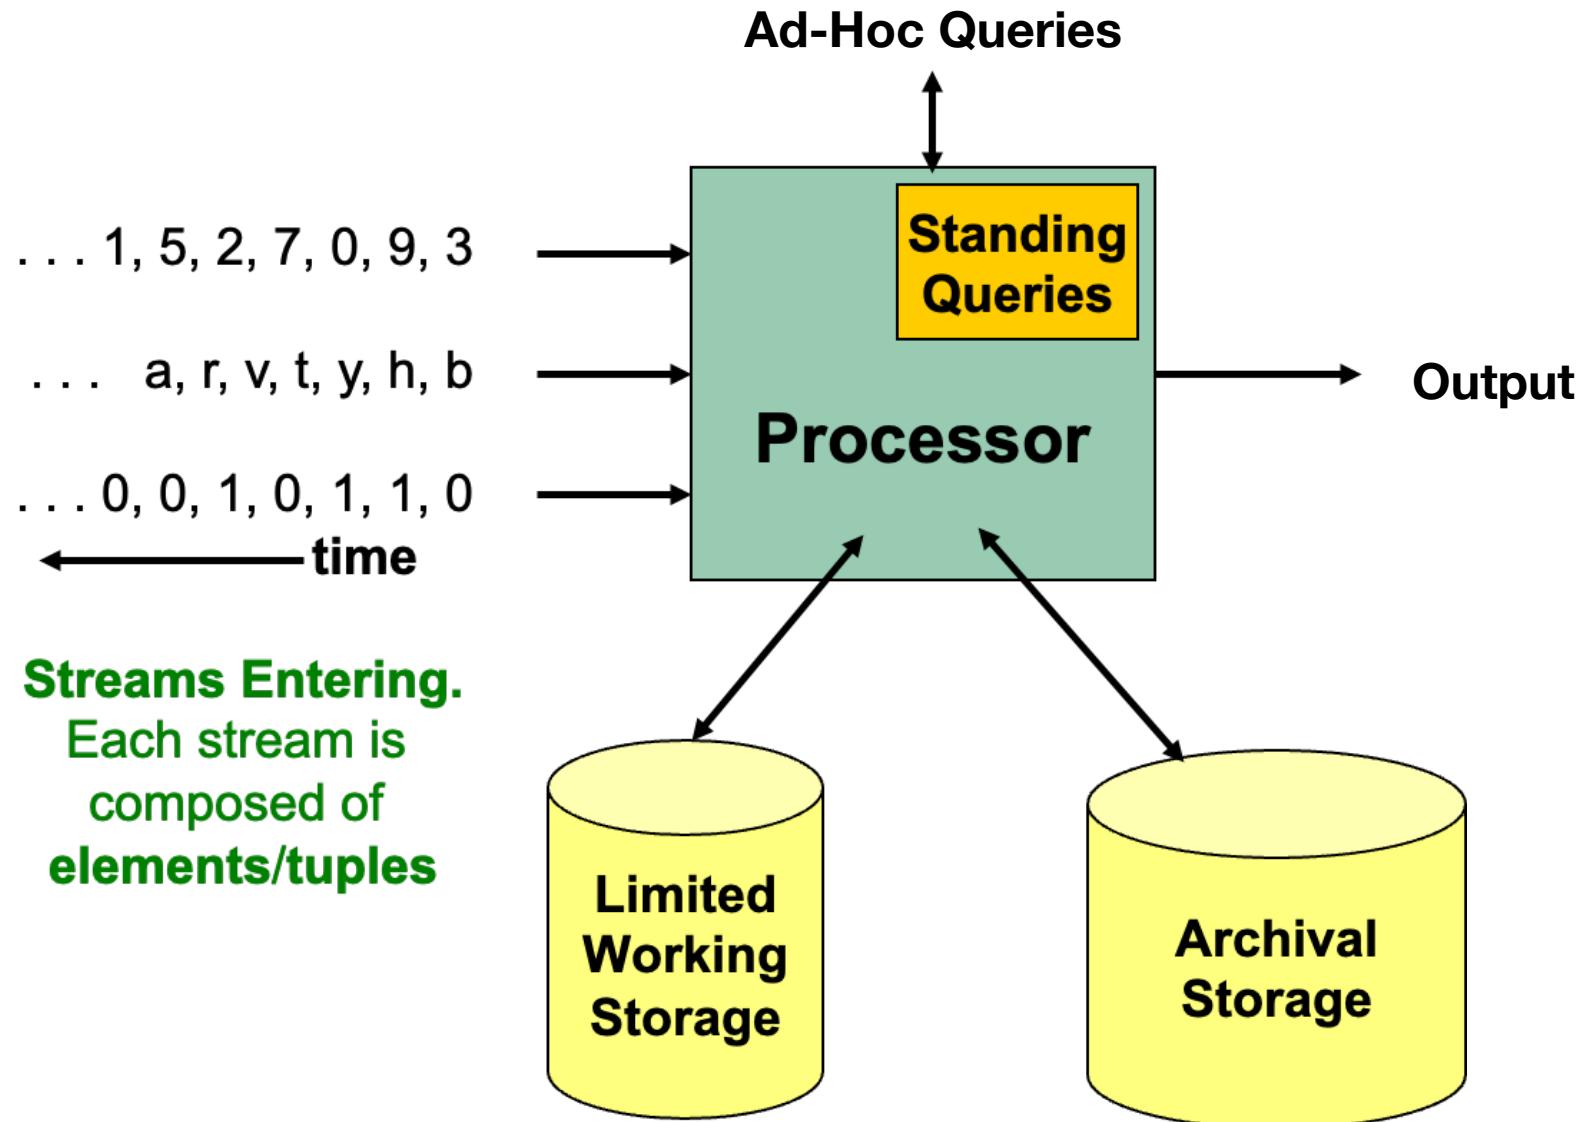


Each time, we only merge buckets of the same color

# Summary

- Sampling a fixed proportion of a stream
  - Sample size grows as the stream grows
- Sampling a fixed-size sample
  - Reservoir sampling
- Counting the number of 1s in the last  $N$  elements
  - Exponentially increasing windows
  - Extensions:
    - Sums of integers in the last  $N$  elements

# Review—General Stream Processing Model



# Review—Sampling from a Data Stream

- Since we **can not store the entire stream**, one obvious approach is to store **a sample**
- How do we sample from the stream just **like we have seen the entire stream?**
  - At any “time”  $k$  we would like a random sample of  $s$  elements
  - **What is the property of the sample we want to maintain?**
    - For all time steps  $k$ , each of  $k$  elements seen so far has equal prob. of being sampled
- Two different problems:
  1. Sample a fixed proportion
  2. Sample a fixed-sized set

# Review – Fixed Size Sample

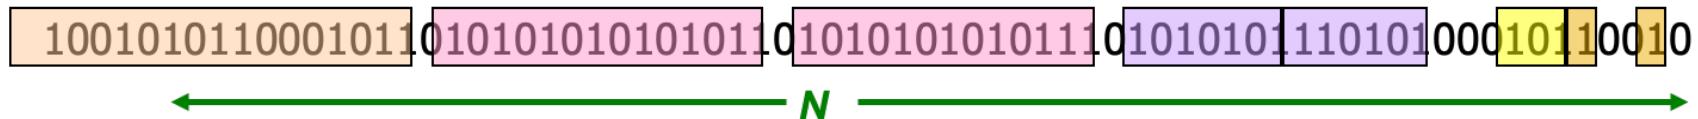
- **Algorithm (a.k.a. Reservoir Sampling)**
  - Store all the first  $s$  elements of the stream to  $\mathbf{S}$
  - Suppose we have seen  $n-1$  elements, and now the  $n^{\text{th}}$  element arrives ( $n > s$ )
    - With probability  $s/n$ , keep the  $n^{\text{th}}$  element, else discard it
    - If we picked the  $n^{\text{th}}$  element, then it replaces one of the  $s$  elements in the sample  $\mathbf{S}$ , picked uniformly at random
- Claim: This algorithm maintains a sample  $\mathbf{S}$  with the desired property:
  - After  $n$  elements, the sample contains each element seen so far with probability  $s/n$

# Review—Sliding Windows

- A useful model of stream processing is that queries are about a **window** of length **N** – the **N** most recent elements received
- **Interesting case:** **N** is so large that the data cannot be stored in memory, or even on disk
  - Or, there are so many streams that windows for all cannot be stored
- Example:
  - For every product **X** we keep 0/1 stream of whether that product was sold in **each** transaction
  - We want to answer queries: how many times have we sold **X** in the last **k** sales (cannot know **k** in advance)

# Review—DGIM: Buckets

- A bucket in the DGIM method is a record consisting of:
  - A. The timestamp of its end expressed in **[ $O(\log N)$  bits]**
  - B. The number of 1s between its beginning and end expressed in **[ $O(\log N - 1) = O(\log N)$  bits]**
- Constraint on buckets: Number of **1s** must be a power of **2**



- How do we filter data streams?

# Filtering Data Streams

- Each element of data stream is a tuple
- Given a list of keys **S**
- Determine which tuples of stream are in **S**
- Obvious solution: **Hash table**
  - But suppose we **do not have enough memory** to store all of **S** in a hash table
    - E.g., we might be processing millions of filters on the same stream

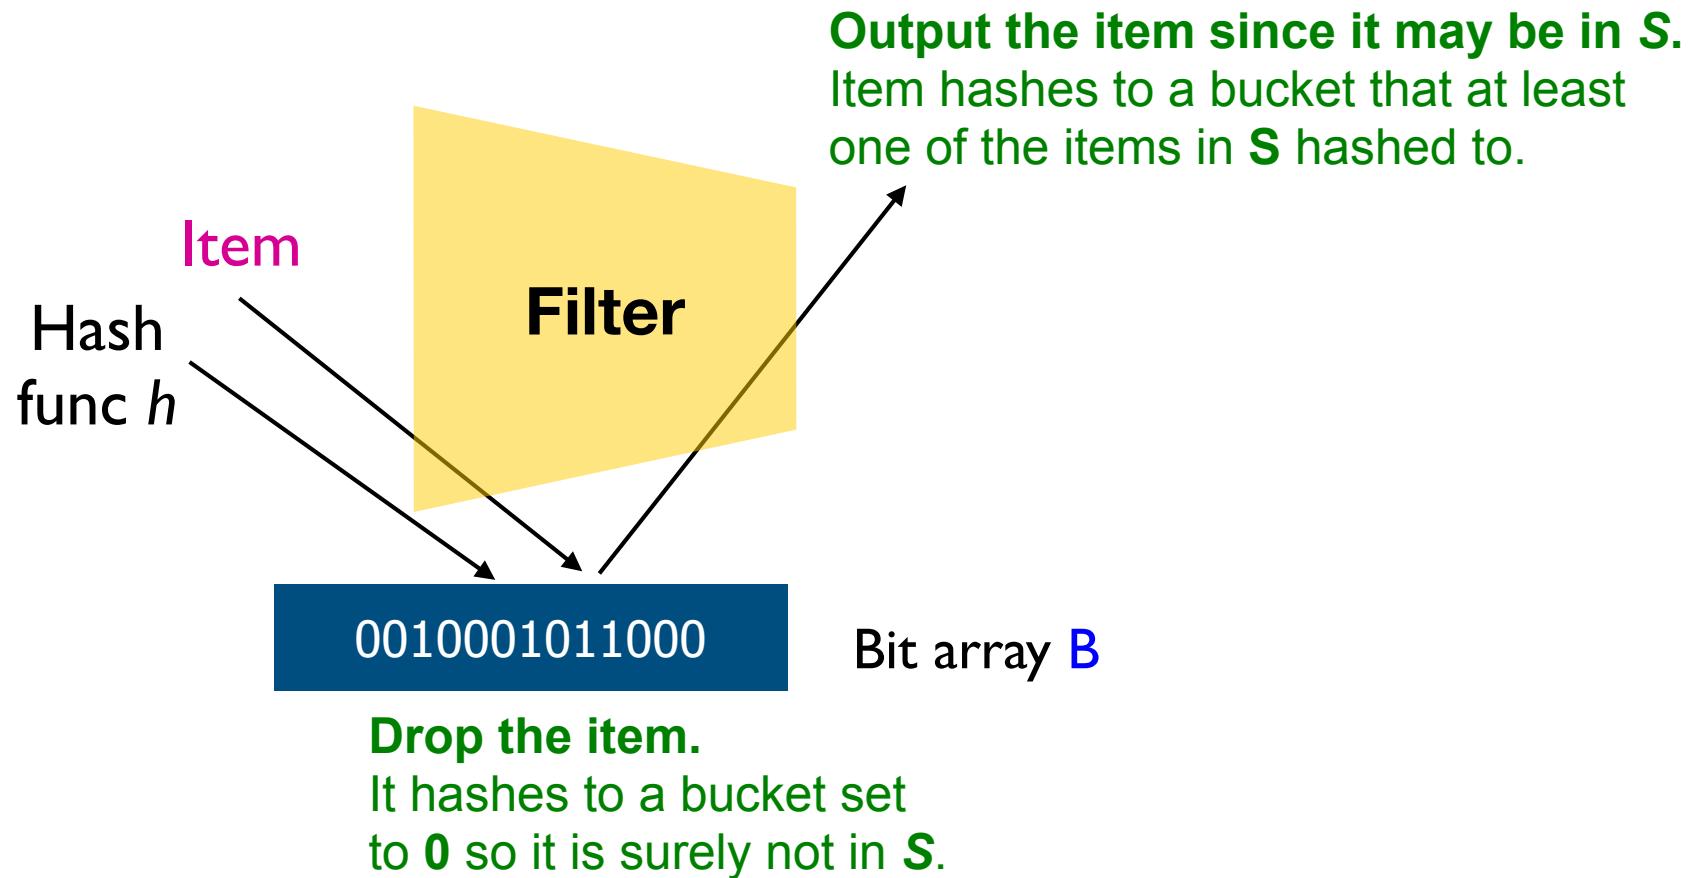
# Applications

- Example: **Email spam filtering**
  - We know 1 billion “good” email addresses
  - If an email comes from one of these, it is **NOT** spam
- Example: **Publish-subscribe systems**
  - You are collecting lots of messages (news articles)
  - People express interest in certain sets of keywords
  - Determine whether each message matches user’s interest

# First Attempt

- Given a set of keys  $S$  that we want to filter
  - Create a **bit array  $B$**  of  $n$  bits, initially all **0s**
  - Choose a **hash function  $h$**  with range  $[0,n)$
  - Hash each member of  $s \in S$  to one of  $n$  buckets, and set that bit to **1**, i.e.,  $B[h(s)] = 1$
  - Hash each element  $a$  of the stream and output only those that hash to bit that was set to **1**
  - Output  $a$  if  $B[h(a)] == 1$

# First Attempt



- Creates **false positives** but no **false negatives**
  - If the item is in  $S$  we surely output it, if not we may still output it

# First Attempt

- $|S| = 1 \text{ billion email addresses}$   
 $|B| = 1 \text{ GB} = 8 \text{ billion bits}$
- If the email address is in **S**, then it surely hashes to a bucket that has the bit set to 1, so it always gets through (**no false negatives**)
- Approximately **1/8** of the bits are set to 1, so about **1/8th** of the addresses not in **S** get through to the output (landing randomly, **false positives**)
  - Actually, less than **1/8th**, because more than one address might hash to the same bit

# Analysis

- More accurate analysis for the number of **false positives**
- Consider: If we throw **m** darts into **n** equally likely targets,  
what is the probability that a target gets at least one dart?
- In our case:
  - Targets = bits or buckets
  - Darts = hash values of items

# Analysis

- We have  $m$  darts,  $n$  targets
- What is the probability that a target gets at least one dart?

The diagram illustrates the derivation of the formula for the probability that a target gets at least one dart. It shows a large rectangle representing the total area of  $n$  targets, divided into  $n$  smaller squares, each representing a target with side length  $m/n$ . The probability that a single dart misses a target is labeled as  $1 - 1/n$ . The probability that all  $n$  targets are missed is  $(1 - 1/n)^n$ . This expression is equivalent to  $e^{-m/n}$ , which is also labeled as  $1 - e^{-m/n}$ . The probability that at least one target is hit is the complement of all targets being missed, which is  $1 - (1 - 1/n)^n$ .

Equals  $1/e$   
as  $n \rightarrow \infty$

Equivalent

$1 - (1 - 1/n)^n$

$n(m/n)$

$1 - e^{-m/n}$

Probability some target  $X$  not hit by a dart

Probability at least one dart hits target  $X$

# Analysis

- Fraction of 1s in the array B (by random throw)=  
= **probability of false positive** =  $1 - e^{-m/n}$
- Example:  $10^9$  darts,  $8 \cdot 10^9$  targets
  - Fraction of **1s** in B =  $1 - e^{-1/8} = 0.1175$
  - Compare with our earlier estimate:  $1/8 = 0.125$

# Bloom Filter

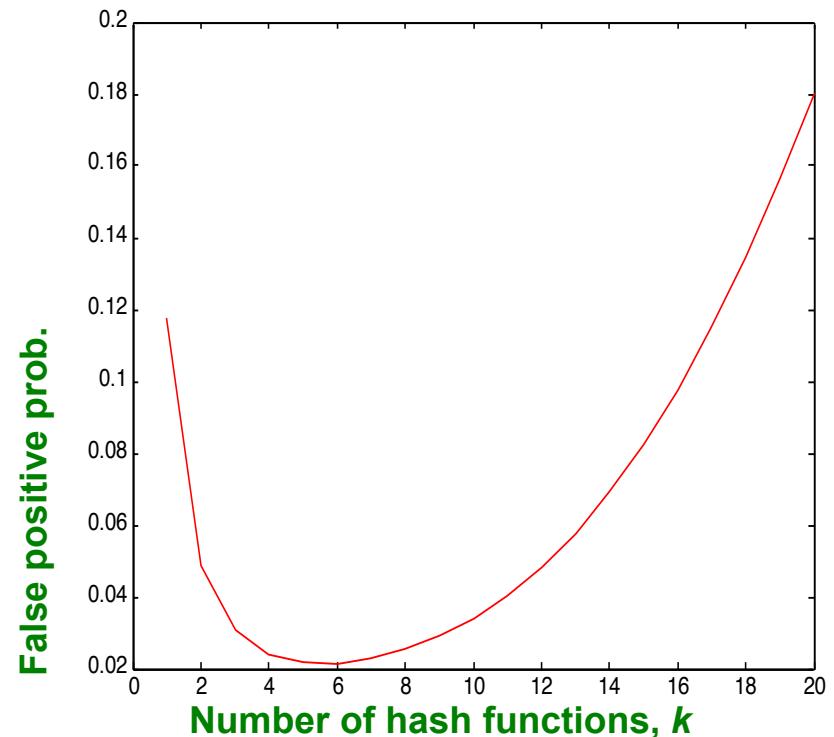
- Consider:  $|S| = m$ ,  $|B| = n$
- Use  $k$  independent hash functions  $h_1, \dots, h_k$
- Initialization:
  - Set  $B$  to all 0s
  - Hash each element  $s \in S$  using each hash function  $h_i$ , set  $B[h_i(s)] = 1$  (for each  $i = 1, \dots, k$ )
- Run-time:
  - When a stream element with key  $x$  arrives
    - If  $B[h_i(x)] = 1$  for all  $i = 1, \dots, k$  then declare that  $x$  is in  $S$
    - That is,  $x$  hashes to a bucket set to 1 for every hash function  $h_i(x)$
  - Otherwise discard the element  $x$

# Bloom Filter – Analysis

- What fraction of the bit vector  $B$  are 1s (by random throw)?
  - Throwing  $k \cdot m$  darts at  $n$  targets
  - So fraction of 1s is  $(1 - e^{-km/n})$
- But we have  $k$  independent hash functions and we only let the element  $x$  through if all  $k$  hash element  $x$  to a bucket of value 1
- So, false positive probability =  $(1 - e^{-km/n})^k$

# Bloom Filter – Analysis

- $m = 1$  billion,  $n = 8$  billion
  - $k = 1$ :  $(1 - e^{-1/8}) = 0.1175$
  - $k = 2$ :  $(1 - e^{-1/4})^2 = 0.0493$
- What happens as we keep increasing  $k$ ?
- “Optimal” value of  $k$ :  $n/m \ln(2)$ 
  - In our case: Optimal  $k = 8 \ln(2) = 5.54 \approx 6$
  - Error at  $k = 6$ :  $(1 - e^{-6/8})^6 = 0.0216$



# Summary

- Bloom filters guarantee no false negatives, and use limited memory
  - Great for pre-processing before more expensive checks
- Suitable for hardware implementation
  - Hash function computations can be parallelized
- Is it better to have **1 big B** or **k small Bs**?
  - It is the same:  $(1 - e^{-km/n})^k$  vs.  $(1 - e^{-m/(n/k)})^k$
  - But keeping **1 big B** is simpler

# Problems on Data Streams

Types of queries one wants to answer on a data stream:

- **Sampling data from a stream**
  - Construct a random sample
- **Queries over sliding windows**
  - Number of items of type **x** in the last **k** elements of the stream
- **Filtering a data stream**
  - Select elements with property **x** from the stream

# Problems on Data Streams

Types of queries one wants to answer on a data stream:

- **Counting distinct elements**
  - Number of distinct elements in the last  $k$  elements of the stream
- **Estimating moments**
  - Estimate avg./std. dev. of last  $k$  elements
- **Finding frequent elements**

- How do we count distinctive items in a stream?

# Counting Distinct Elements

- Problem:
  - Data stream consists of a universe of elements chosen from a set of size  $N$
  - Maintain a count of the number of distinct elements seen so far
- Obvious approach:  
Maintain the set of elements seen so far
  - That is, keep a hash table of all the distinct elements seen so far

# Applications

- How many different words are found among the Web pages being crawled at a site?
  - Unusually low or high numbers could indicate artificial pages (spam?)
- How many different web pages does each customer request in a week?
- How many distinct products have we sold in the last week?

# Using Small Storage

- Real problem: What if we do not have space to maintain the set of elements seen so far?
  - Estimate the count in an unbiased way
  - Accept that the count may have a little error, but limit the probability that the error is large

# Flajolet-Martin Approach

- Pick a hash function  $h$  that maps each of the  $N$  elements to at least  $\log_2 N$  bits
- For each stream element  $a$ , let  $r(a)$  be the number of trailing 0s in  $h(a)$ 
  - $r(a) = \text{position of first 1 counting from the right}$ 
    - E.g., say  $h(a) = 12$ , then 12 is 1100 in binary, so  $r(a) = 2$
- Record  $R = \text{the maximum } r(a) \text{ seen}$ 
  - $R = \max_a r(a)$ , over all the items  $a$  seen so far
- Estimated number of distinct elements =  $2^R$

# Intuition

- Very rough and heuristic intuition why Flajolet-Martin works:
  - $h(a)$  hashes  $a$  with **equal prob.** to any of  $N$  values
  - Then  $h(a)$  is a sequence of  $\log_2 N$  bits, where  $2^{-r}$  fraction of all  $as$  have a tail of  $r$  zeros
    - About 50% of  $as$  hash to \*\*\*0
    - About 25% of  $as$  hash to \*\*00
    - So, if we saw the longest tail of  $r=2$  (i.e., item hash ending \*100) then we have probably seen **about 4** distinct items so far
  - So, it takes to hash about  $2^r$  items before we see one with zero-suffix of length  $r$

# Formal Derivation

- Now we show why Flajolet-Martin works
- Formally, letting  $m$  be the number of distinct elements seen so far in the stream, we will show that probability of finding a tail of  $r$  zeros:
  - Goes to 1 if  $m \gg 2^r$
  - Goes to 0 if  $m \ll 2^r$
- Thus,  $2^R$  will almost always be around  $m$

# Formal Derivation

- The probability that a given  $h(a)$  ends in at least  $r$  zeros is  $2^{-r}$ 
  - $h(a)$  hashes elements uniformly at random
  - Probability that a random number ends in at least  $r$  zeros is  $2^{-r}$
- Then, the probability of **NOT** seeing a tail of length  $r$  among  $m$  elements:

$$(1 - 2^{-r})^m$$

Prob. all end in fewer than  $r$  zeros.

Prob. that given  $h(a)$  ends in fewer than  $r$  zeros

# Formal Derivation

- Note:  $(1 - 2^{-r})^m = (1 - 2^{-r})^{2^r(m2^{-r})} \approx e^{-m2^{-r}}$
- Prob. of NOT finding a tail of length r is:
  - If  $m \ll 2^r$ , then prob. tends to 1
    - $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 1$  as  $m/2^r \rightarrow 0$
    - So, the probability of finding a tail of length r tends to 0
  - If  $m \gg 2^r$ , then prob. tends to 0
    - $(1 - 2^{-r})^m \approx e^{-m2^{-r}} = 0$  as  $m/2^r \rightarrow \infty$
    - So, the probability of finding a tail of length r tends to 1
- Thus,  $2^R$  will almost always be around m

# Fixups

- $E[2^R]$  is actually infinite
  - Probability of seeking an item with trailing 0s halves when  $R \rightarrow R+1$ , but estimated number of distinct items doubles
- Workaround involves using many hash functions  $h_i$  and getting many samples of  $R_i$ ,
- How are samples  $R_i$  combined?
  - Average? What if one very large value  $2^{R_i}$ ?
  - Median? All estimates are a power of 2
- Solution:
  - Partition your samples into small groups
  - Take the average of groups
  - Then take the median of the averages

- How do we calculate statistics for the last  $k$  items in a stream?

# Generalization: Moments

- Suppose a stream has elements chosen from a set **A** of **N** values
- Let  $m_i$  be the number of times value **i** occurs in the stream
- The  $k^{\text{th}}$  moment is

$$\sum_{i \in A} (m_i)^k$$

# Cases

$$\sum_{i \in A} (m_i)^k$$

- **0<sup>th</sup>moment** = number of distinct elements
  - The problem just considered
- **1<sup>st</sup> moment** = count of the numbers of elements = length of the stream
  - Easy to compute
- **2<sup>nd</sup> moment** = **surprise number S** = a measure of how uneven the distribution is

# Example: Surprise Number

- Stream of length 100
- 11 distinct values
- Item counts: 10, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9
  - **Surprise S = 910**
- Item counts: 90, 1, 1, 1, 1, 1, 1, 1 ,1, 1, 1
  - **Surprise S = 8110**

# AMS Method

- AMS method works for all moments
- Gives an unbiased estimate
- We will just concentrate on the **2<sup>nd</sup> moment S**
- We pick and keep track of many variables **X**:
  - For each variable **X** we store **X.el** and **X.val**
    - **X.el** corresponds to the item **i**
    - **X.val** corresponds to the **count** of item **i**
  - Note this requires a count in main memory, so number of **Xs** is limited
- Our goal is to compute 
$$S = \sum_i m_i^2$$

# One Random Variable ( $X$ )

- How to set  $X.val$  and  $X.el$ ?
  - Assume stream has length  $n$  (we relax this later)
  - Pick some random time  $t$  ( $t < n$ ) to start, so that **any time is equally likely**, and we pick  $k$  of it
  - Let at time  $t$  the stream have item  $i$ . **We set  $X.el = i$**
  - Then we maintain count  $c(i)$  ( **$X.val = c$** ) of the number of item  **$i$**  is in the stream starting from the chosen time  $t$
- Then the estimate of the **2<sup>nd</sup> moment**  $S = \sum_i m_i^2$  is:

$$f(X_i) = n(2c(i) - 1), \quad S = \frac{1}{k} \sum_i^k f(X_i)$$

# An Example

- Suppose the stream is **a, b, c, b, d, a, c, d, a, b, d, c, a, a, b.**  
The length of the stream is **n = 15**
- Since **a** appears **5** times, **b** appears **4** times, **c** appears **3** times, and **d** appears **3** times, the second moment for the stream is  $5^2 + 4^2 + 3^2 + 3^2 = 59$
- We keep three variables **X1, X2, and X3** and pick at random the 3rd, 8th, and 13th positions to define them

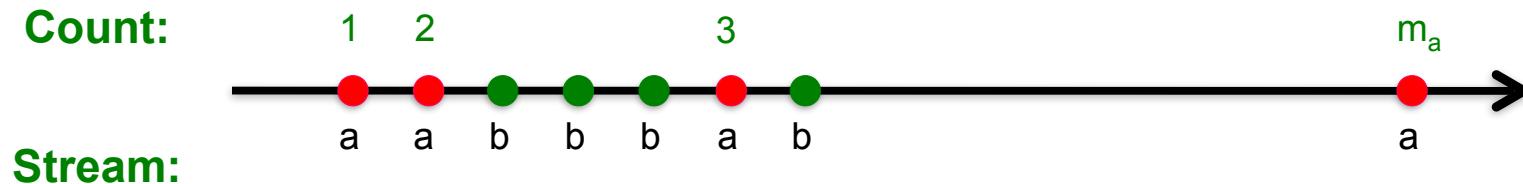
**X1.element = c, X1.value = 3** From X1, we derive  $n(2X1.value - 1) = 15 \times (2 \times 3 - 1) = 75$

**X2.element = d, X2.value = 2** From X2, we derive  $n(2X2.value - 1) = 15 \times (2 \times 2 - 1) = 45$

**X3.element = a, X3.value = 2** From X3, we derive  $n(2X3.value - 1) = 15 \times (2 \times 2 - 1) = 45$

- The average of three estimates is 55, which is close to 59

# Expectation Analysis

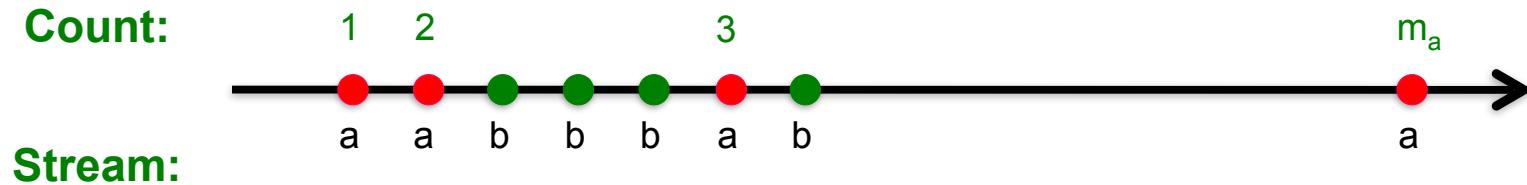


- 2<sup>nd</sup> moment is  $S = \sum_i m_i^2$
- $c_t$  ... number of times item at time  $t$  appears from time  $t$  onwards ( $c_1=m_a$ ,  $c_2=m_a-1$ ,  $c_3=m_b$ )
- The expected value of  $n(2c(i) - 1)$  is the average over all positions  $i$  between 1 and  $n$  of  $n(2c(i) - 1)$ , that is

$$E[f(X)] = \frac{1}{n} \sum_{i=1}^n n(2c(i) - 1) = \frac{1}{n} \sum_n n (1 + 3 + 5 + \dots + 2m_a - 1)$$

Group times by the value seen      Time  $t$  when the last  $X$  is seen ( $c(i)=1$ )      Time  $t$  when the penultimate  $X$  is seen ( $c(i)=2$ )      Time  $t$  when the first  $X$  is seen ( $c(i)=m_a$ )

# Expectation Analysis



- $E[f(X)] = \frac{1}{n} \sum_i n (1 + 3 + 5 + \dots + 2m_i - 1)$ 
  - $(1 + 3 + 5 + \dots + 2m_i - 1) = \sum_{i=1}^{m_i} (2i - 1) = 2 \frac{m_i(m_i + 1)}{2} - m_i = (m_i)^2$
  - Then  $E[f(X)] = \frac{1}{n} \sum_i n (m_i)^2$
  - So,  $\mathbf{E}[f(X)] = \sum_i (m_i)^2 = S$
- We have the second moment (in expectation)!

# Higher-Order Moments

- For estimating  $k^{\text{th}}$  moment we essentially use the same algorithm but change the estimate:
  - For  $k=2$  we used  $n (2 \cdot c - 1)$
  - For  $k=3$  we use:  $n (3 \cdot c^2 - 3c + 1)$  (where  $c=X.\text{val}$ )
- Why?
  - For  $k=2$ : Remember we had  $(1 + 3 + 5 + \dots + 2m_a - 1)$  and we showed terms  $2c-1$  (for  $c=1, \dots, m$ ) sum to  $m^2$ 
    - $\sum_{c=1}^m 2c - 1 = \sum_{c=1}^m c^2 - \sum_{c=1}^m (c - 1)^2 = m^2$
    - So:  $2c - 1 = c^2 - (c - 1)^2$
  - For  $k=3$ :  $c^3 - (c-1)^3 = 3c^2 - 3c + 1$
  - Generally: Estimate =  $n (c^k - (c - 1)^k)$

# Combining Samples

- In practice:
  - Compute  $f(\mathbf{X}) = n(2c - 1)$  for as many variables  $\mathbf{X}$  as you can fit in memory
  - Average them in groups
  - Take median of averages
- Problem: Streams never end
  - We assumed there was a number  $n$ , the number of positions in the stream
  - But real streams go on forever, so  $n$  is a variable – the number of inputs seen so far

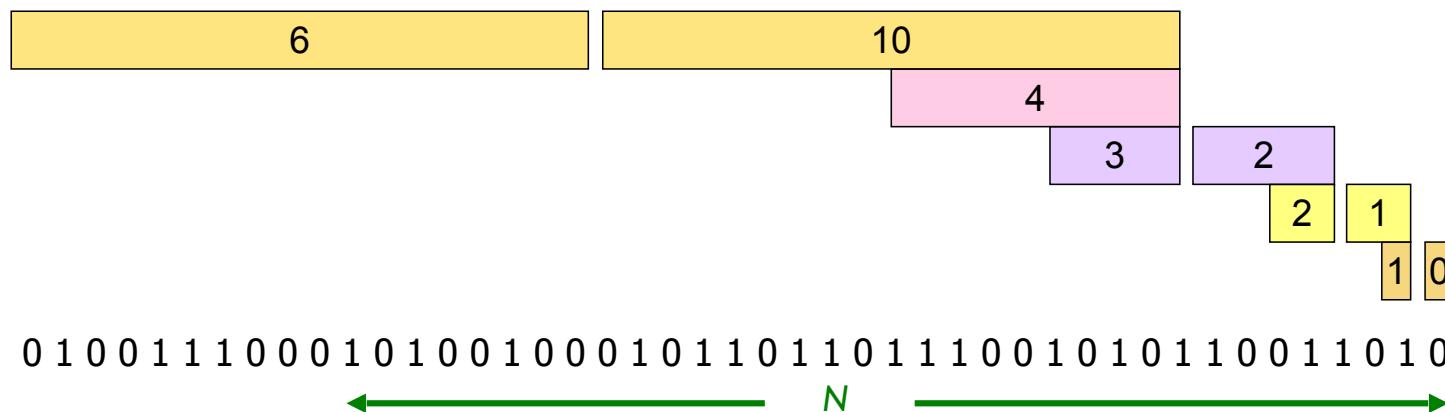
# Streams Never End: Fixups

- (1) The variables  $\mathbf{X}$  have  $n$  as a factor – keep  $n$  separately; just hold the count in  $\mathbf{X}$
- (2) Suppose we can only store  $k$  counts. We must throw some  $\mathbf{X}$ s out as time goes on:
  - **Objective:** Each starting time  $t$  is selected with probability  $k/n$
  - **Solution:** (fixed-size sampling!)
    - Choose the first  $k$  times for  $k$  variables
    - When the  $n^{\text{th}}$  element arrives ( $n > k$ ), choose it with probability  $k/n$
    - If you choose it, throw one of the previously stored variables  $\mathbf{X}$  out, with equal probability

- How do we count frequently-appearing items in a stream?

# Counting Itemsets

- **New Problem:** Given a stream, which items appear more than  $s$  times in the window?
- **Possible solution:** Tear the stream up into multiple binary streams; one binary stream per item
  - 1 = item present; 0 = not present
  - Use **DGIM** to estimate counts of 1s for all items



# Extensions

- In principle, you could count frequent pairs or even larger sets the same way
  - One stream per itemset
- Drawbacks:
  - Only approximate
  - **Number of itemsets** is way too big

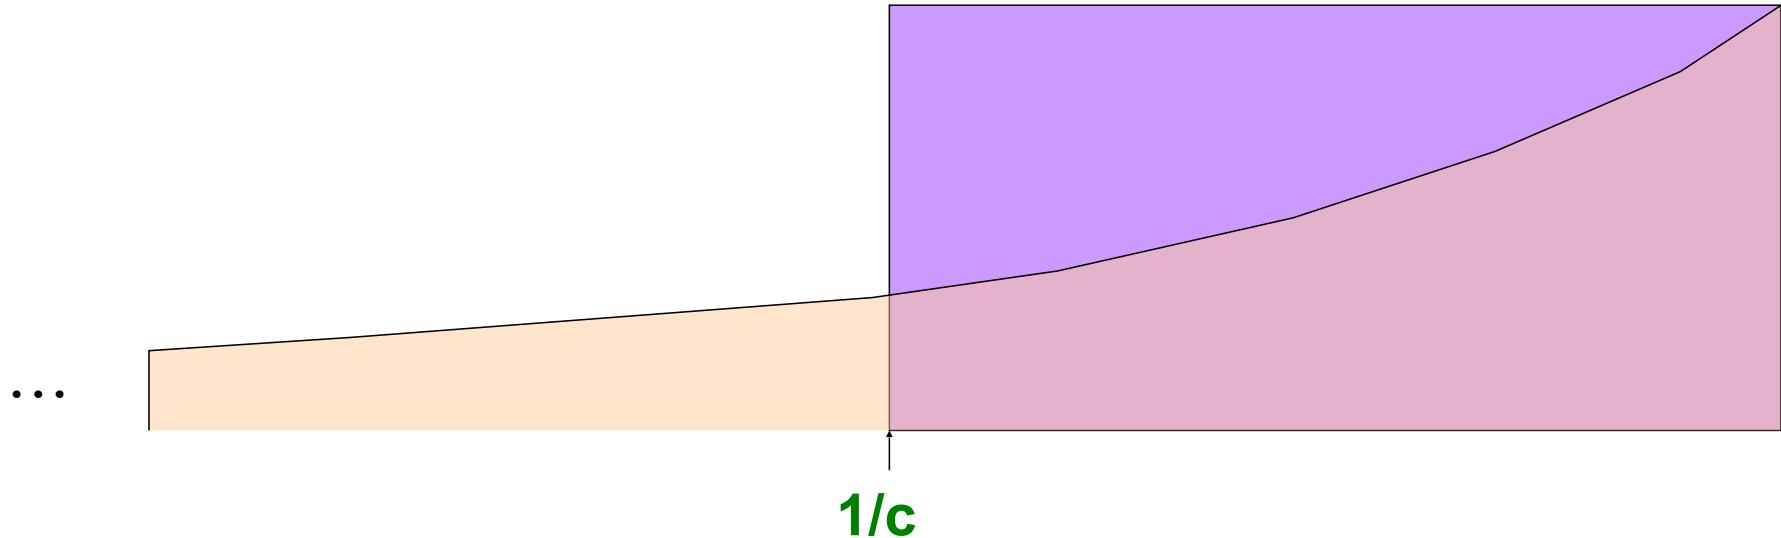
# Exponentially Decaying Windows

- Exponentially decaying windows: A heuristic for selecting likely frequent item(sets)
  - What are “currently” most popular movies?
    - Instead of computing the raw count in last **N** elements
    - Compute a **smooth aggregation** over the whole stream
  - If stream is  $a_1, a_2, \dots$  and we are taking the sum of the stream, take the answer at time **t** to be:  $= \sum_{i=1}^t a_i (1 - c)^{t-i}$ 
    - **c** is a constant, presumably tiny, like  $10^{-6}$  or  $10^{-9}$
    - When **new  $a_{t+1}$**  arrives:  
Multiply current sum by **(1-c)** and add  $a_{t+1}$

# Example: Counting Items

- If each  $a_i$  is an “item” we can compute the characteristic function of each possible item  $x$  as an Exponentially Decaying Window
  - That is:  $\sum_{i=1}^t \delta_i \cdot (1 - c)^{t-i}$  where  $\delta_i = 1$  if  $a_i = x$ , and  $0$  otherwise
  - Imagine that for each item  $x$  we have a binary stream (1 if  $x$  appears, 0 if  $x$  does not appear)
  - New item  $x$  arrives:
    - Multiply all counts by  $(1 - c)$
    - Add +1 to count for element  $x$
  - Call this sum the “weight” of item  $x$

# Sliding Versus Decaying Windows



- Important property: Sum over all weights  $\sum_t (1 - c)^t$  is

$$1/[1 - (1 - c)] = 1/c$$

$$\sum_{k=0}^n z^k = \frac{1 - z^{n+1}}{1 - z}$$

# Example: Counting Items

- What are “currently” most popular movies?
- Suppose we want to find movies of weight  $> \frac{1}{2}$ 
  - Important property: Sum over all weights  $\sum_t (1 - c)^t$  is
$$\frac{1}{1 - (1 - c)} = \frac{1}{c}$$
- Thus:
  - There cannot be more than  $2/c$  movies with weight of  $\frac{1}{2}$  or more
  - So,  $2/c$  is a **limit** on the **number of movies** being counted at any time

# Reading

- Jure Leskovec, Anand Raj, Jeff Ullman, “Mining of Massive Datasets,” Cambridge University Press, Chapter 4